

classifying-penguins

February 13, 2025

1 Classifying Penguins with Decision Trees

Load the penguins dataset - as before, we'll use Seaborn dataset:

```
[1]: import seaborn as sns
import pandas as pd
import matplotlib.pyplot as plt
```

```
[2]: penguins_df = sns.load_dataset("penguins")
penguins_df.head()
```

```
[2]: species      island  bill_length_mm  bill_depth_mm  flipper_length_mm  \
0  Adelie  Torgersen         39.1           18.7           181.0
1  Adelie  Torgersen         39.5           17.4           186.0
2  Adelie  Torgersen         40.3           18.0           195.0
3  Adelie  Torgersen          NaN           NaN           NaN
4  Adelie  Torgersen         36.7           19.3           193.0

      body_mass_g      sex
0         3750.0    Male
1         3800.0  Female
2         3250.0  Female
3            NaN     NaN
4         3450.0  Female
```

```
[3]: len(penguins_df)
```

```
[3]: 344
```

```
[4]: penguins_df.describe().T
```

```
[4]:
```

	count	mean	std	min	25%	50%	\
bill_length_mm	342.0	43.921930	5.459584	32.1	39.225	44.45	
bill_depth_mm	342.0	17.151170	1.974793	13.1	15.600	17.30	
flipper_length_mm	342.0	200.915205	14.061714	172.0	190.000	197.00	
body_mass_g	342.0	4201.754386	801.954536	2700.0	3550.000	4050.00	
		75%	max				

```
bill_length_mm      48.5    59.6
bill_depth_mm       18.7    21.5
flipper_length_mm   213.0   231.0
body_mass_g         4750.0  6300.0
```

```
[5]: penguins_df.isnull().sum()
```

```
[5]: species          0
     island          0
     bill_length_mm    2
     bill_depth_mm    2
     flipper_length_mm 2
     body_mass_g       2
     sex              11
     dtype: int64
```

We've spotted some missing values. For this example we'll just get rid of rows with missing values.

```
[6]: penguins_df = penguins_df.dropna()
```

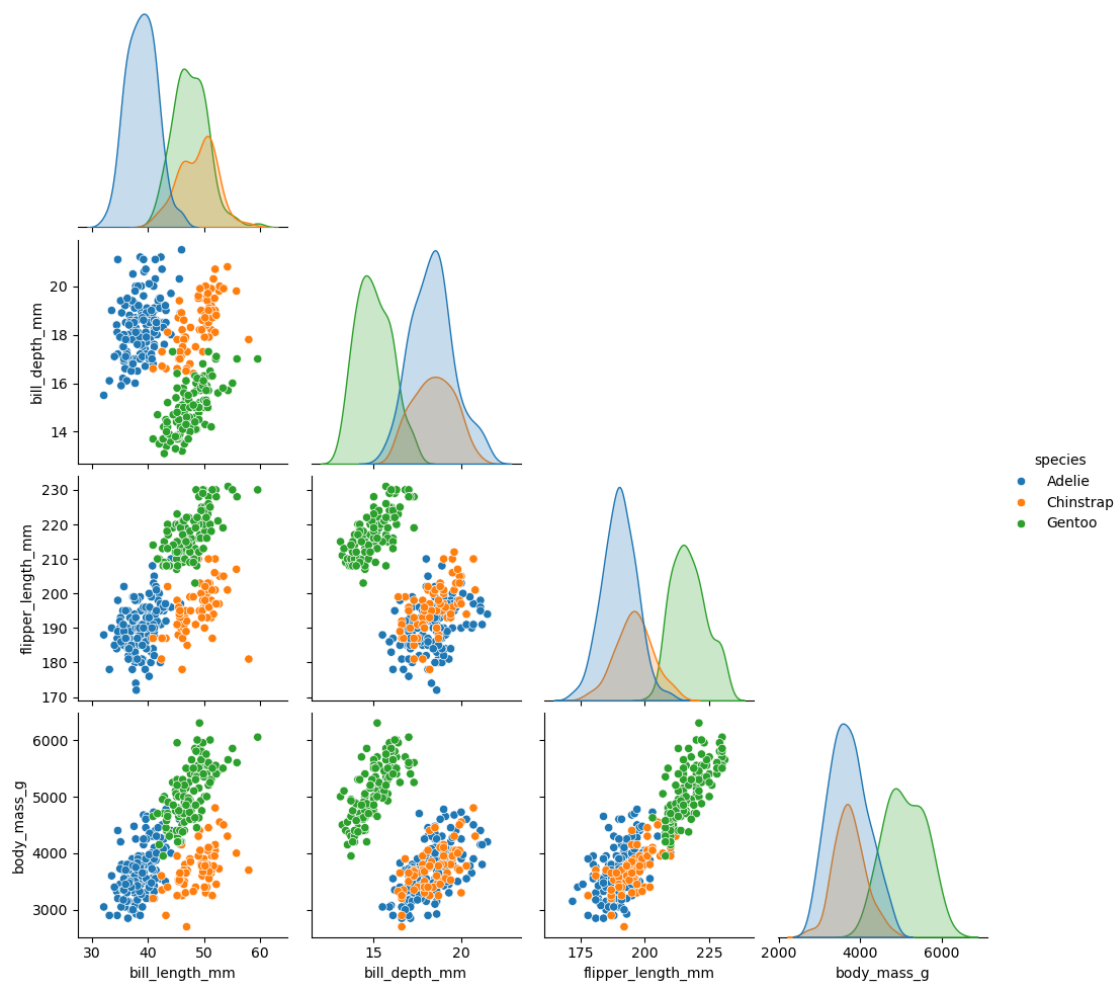
```
[7]: len(penguins_df)
```

```
[7]: 333
```

Let's get a quick visual on the relationship between all pairs of numeric columns.

```
[8]: sns.pairplot(data=penguins_df,
                  hue="species",
                  corner=True)
```

```
[8]: <seaborn.axisgrid.PairGrid at 0x2bb88d25eb0>
```



1.1 Using a Decision Tree Classifier to Predict Penguin Species

```
[9]: penguins_df.head()
```

```
[9]:  species    island  bill_length_mm  bill_depth_mm  flipper_length_mm  \
0  Adelie  Torgersen         39.1           18.7           181.0
1  Adelie  Torgersen         39.5           17.4           186.0
2  Adelie  Torgersen         40.3           18.0           195.0
4  Adelie  Torgersen         36.7           19.3           193.0
5  Adelie  Torgersen         39.3           20.6           190.0

    body_mass_g    sex
0       3750.0  Male
1       3800.0 Female
2       3250.0 Female
4       3450.0 Female
```

5 3650.0 Male

1.1.1 Encode categorical variables

The species, island, and sex variables are strings. This will cause a problem for our decision tree classifier which **works with numeric variables**. To solve this, we'll need to encode those variables.

See: <https://scikit-learn.org/stable/modules/generated/sklearn.preprocessing.LabelEncoder.html#sklearn.preprocessing.LabelEncoder>
https://scikit-learn.org/stable/modules/preprocessing_targets.html and <https://scikit-learn.org/stable/modules/generated/sklearn.preprocessing.OrdinalEncoder.html#sklearn.preprocessing.OrdinalEncoder>

The species is our label, so we'll use a LabelEncoder for that.

```
[10]: from sklearn.preprocessing import LabelEncoder

label_encoder = LabelEncoder()
```

Let's first make a quick DataFrame with two columns: the original penguin species and it's new encoded value.

```
[11]: encoding_df = pd.DataFrame({"species": penguins_df["species"],
                                "species_encoded": label_encoder.
                                ↪fit_transform(penguins_df["species"])}))
```

```
[12]: encoding_df.sample(10)
```

```
[12]:
```

	species	species_encoded
190	Chinstrap	1
172	Chinstrap	1
195	Chinstrap	1
34	Adelie	0
78	Adelie	0
208	Chinstrap	1
156	Chinstrap	1
293	Gentoo	2
119	Adelie	0
95	Adelie	0

Now we'll replace the original text based penguin species label with it's numerically encoded version which will be compatible with our machine learning models.

```
[13]: penguins_df['species'] = label_encoder.fit_transform(penguins_df['species'])
```

We can convert the male/female values in the sex column using an OrdinalEncoder. This will change them to 0 and 1.

```
[14]: from sklearn.preprocessing import OrdinalEncoder

ordinal_encoder = OrdinalEncoder()
```

Again, let's make a quick DataFrame to show the encoded values.

```
[15]: encoded_df = pd.DataFrame({"sex": penguins_df["sex"]})
      encoded_df[['encoded_value']] = ordinal_encoder.
          ↪ fit_transform(encoded_df[['sex']])

      encoded_df.head()
```

```
[15]:      sex  encoded_value
0   Male             1.0
1  Female             0.0
2  Female             0.0
4  Female             0.0
5   Male             1.0
```

Now we'll go ahead and update the penguins sex column with the numerically encoded values.

```
[16]: penguins_df[['sex']] = ordinal_encoder.fit_transform(penguins_df[['sex']])
```

```
[17]: penguins_df.head()
```

```
[17]:   species  island  bill_length_mm  bill_depth_mm  flipper_length_mm  \
0         0  Torgersen           39.1           18.7           181.0
1         0  Torgersen           39.5           17.4           186.0
2         0  Torgersen           40.3           18.0           195.0
4         0  Torgersen           36.7           19.3           193.0
5         0  Torgersen           39.3           20.6           190.0

      body_mass_g  sex
0         3750.0  1.0
1         3800.0  0.0
2         3250.0  0.0
4         3450.0  0.0
5         3650.0  1.0
```

It makes more sense to use a OneHotEncoder for our island column. As we'll see, this will make our decision tree more readable as we'll avoid questions like "Is the island less than 2?". Instead, we'll get three columns for island with a 1 in the column that corresponds to the original island value, and 0s for the other two.

```
[18]: from sklearn.preprocessing import OneHotEncoder
```

```
[19]: one_hot_encoder = OneHotEncoder()

      island_encoded = one_hot_encoder.fit_transform(penguins_df[['island']])
```

The OneHotEncoder has given us a sparse matrix.

```
[20]: island_encoded
```

```
[20]: <333x3 sparse matrix of type '<class 'numpy.float64'>'
      with 333 stored elements in Compressed Sparse Row format>
```

Before converting this to a Pandas DataFrame, we need to convert it to a dense array.

```
[21]: island_encoded = island_encoded.toarray()

island_encoded
```

[illegible]

7

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

```
[1., 0., 0.],
[1., 0., 0.],
[1., 0., 0.],
[1., 0., 0.],
[1., 0., 0.],
[1., 0., 0.],
[1., 0., 0.],
[1., 0., 0.],
[1., 0., 0.],
[1., 0., 0.],
[1., 0., 0.],
[1., 0., 0.]]
```

Now we can convert the one-hot encoded 'island' result to a DataFrame.

```
[22]: island_encoded_df = pd.DataFrame(
        island_encoded, columns=one_hot_encoder.get_feature_names_out(['island'])
    )

island_encoded_df
```

```
[22]:
```

	island_Biscoe	island_Dream	island_Torgersen
0	0.0	0.0	1.0
1	0.0	0.0	1.0
2	0.0	0.0	1.0
3	0.0	0.0	1.0
4	0.0	0.0	1.0
..
328	1.0	0.0	0.0
329	1.0	0.0	0.0
330	1.0	0.0	0.0
331	1.0	0.0	0.0
332	1.0	0.0	0.0

```
[333 rows x 3 columns]
```

We now bring the new one-hot encoded 'island' columns into the original DataFrame and drop the redundant 'island' column.

```
[23]: penguins_df = pd.concat([penguins_df, island_encoded_df],
                               axis=1, # axis=1 for columns
                               join='inner') # without this we'll introduce NaNs
```

```
[24]: penguins_df = penguins_df.drop(columns=['island'])
```

Now when we check out our DataFrame, we'll see that species, island, and sex are now encoded with numeric values.

```
[25]: penguins_df.sample(15)
```

```
[25]:
```

	species	bill_length_mm	bill_depth_mm	flipper_length_mm	body_mass_g	\
209	1	49.3	19.9	203.0	4050.0	
99	0	43.2	18.5	192.0	4100.0	
309	2	52.1	17.0	230.0	5550.0	
144	0	37.3	16.8	192.0	3000.0	
27	0	40.5	17.9	187.0	3200.0	
295	2	48.6	16.0	230.0	5800.0	
136	0	35.6	17.5	191.0	3175.0	
20	0	37.8	18.3	174.0	3400.0	
259	2	48.7	15.7	208.0	5350.0	
254	2	49.1	14.8	220.0	5150.0	
306	2	43.4	14.4	218.0	4600.0	
202	1	48.1	16.4	199.0	3325.0	
325	2	46.8	16.1	215.0	5500.0	
284	2	45.8	14.2	219.0	4700.0	
56	0	39.0	17.5	186.0	3550.0	

	sex	island_Biscoe	island_Dream	island_Torgersen
209	1.0	0.0	1.0	0.0
99	1.0	1.0	0.0	0.0
309	1.0	1.0	0.0	0.0
144	0.0	0.0	1.0	0.0
27	0.0	0.0	1.0	0.0
295	1.0	1.0	0.0	0.0
136	0.0	0.0	1.0	0.0
20	0.0	1.0	0.0	0.0
259	1.0	1.0	0.0	0.0
254	0.0	1.0	0.0	0.0
306	0.0	1.0	0.0	0.0
202	0.0	0.0	1.0	0.0
325	1.0	1.0	0.0	0.0
284	0.0	1.0	0.0	0.0
56	0.0	1.0	0.0	0.0

1.2 Training and test data

1.2.1 Define X and y

We want to be able to predict a penguin’s species from the observations about that penguin - the “features”. Our target variable is species, which is conventionally called y. The features are conventionally called X (upper case).

```
[26]: X = penguins_df.drop('species', axis=1)
      y = penguins_df['species']
```

```
[27]: X.head()
```

```
[27]:
```

	bill_length_mm	bill_depth_mm	flipper_length_mm	body_mass_g	sex	\
0	39.1	18.7	181.0	3750.0	1.0	
1	39.5	17.4	186.0	3800.0	0.0	
2	40.3	18.0	195.0	3250.0	0.0	
4	36.7	19.3	193.0	3450.0	0.0	
5	39.3	20.6	190.0	3650.0	1.0	

	island_Biscoe	island_Dream	island_Torgersen
0	0.0	0.0	1.0
1	0.0	0.0	1.0
2	0.0	0.0	1.0
4	0.0	0.0	1.0
5	0.0	0.0	1.0

```
[28]: y.head()
```

```
[28]: 0    0
      1    0
      2    0
      4    0
      5    0
      Name: species, dtype: int32
```

1.2.2 Split the dataset into training and testing sets

We are going to build our decision tree using the penguin data, but we'll set aside some of that data so that we can test how accurate our decision tree is on examples that weren't used in its construction. For this we can use the `train_test_split` function.

```
[29]: from sklearn.model_selection import train_test_split
```

```
[30]: X_train, X_test, y_train, y_test = train_test_split(
      X,
      y,
      test_size=0.3,
      random_state=42)
```

```
[31]: X_train.head()
```

```
[31]:
```

	bill_length_mm	bill_depth_mm	flipper_length_mm	body_mass_g	sex	\
268	44.9	13.3	213.0	5100.0	0.0	
87	36.9	18.6	189.0	3500.0	0.0	
154	51.3	19.2	193.0	3650.0	1.0	
44	37.0	16.9	185.0	3000.0	0.0	
229	46.8	15.4	215.0	5150.0	1.0	

	island_Biscoe	island_Dream	island_Torgersen
268	1.0	0.0	0.0

87	0.0	1.0	0.0
154	0.0	1.0	0.0
44	1.0	0.0	0.0
229	1.0	0.0	0.0

```
[32]: y_train.head()
```

```
[32]: 268    2
      87     0
      154    1
      44     0
      229    2
      Name: species, dtype: int32
```

```
[33]: y_train.isnull().sum()
```

```
[33]: 0
```

```
[34]: X_test.head()
```

```
[34]:      bill_length_mm  bill_depth_mm  flipper_length_mm  body_mass_g  sex  \
138           37.0         16.5           185.0        3400.0  0.0
114           39.6         20.7           191.0        3900.0  0.0
143           40.7         17.0           190.0        3725.0  1.0
14            34.6         21.1           198.0        4400.0  1.0
186           49.7         18.6           195.0        3600.0  1.0

      island_Biscoe  island_Dream  island_Torgersen
138             0.0             1.0              0.0
114             0.0             0.0              1.0
143             0.0             1.0              0.0
14              0.0             0.0              1.0
186             0.0             1.0              0.0
```

```
[35]: y_test.head()
```

```
[35]: 138    0
      114    0
      143    0
      14     0
      186    1
      Name: species, dtype: int32
```

1.3 Create our decision tree

Scikit-learn makes this very easy for us with the `DecisionTreeClassifier` and its `fit` method.

```
[36]: from sklearn.tree import DecisionTreeClassifier
```



```
[37]: clf = DecisionTreeClassifier()
      clf = clf.fit(X_train, y_train)
```

Now we use the classifier to predict species in the test set and calculate accuracy

```
[38]: y_pred = clf.predict(X_test)
      y_pred
```

```
[38]: array([0, 0, 0, 0, 1, 0, 2, 1, 2, 0, 1, 1, 2, 0, 1, 2, 2, 0, 0, 0, 2, 2,
           0, 0, 0, 1, 1, 0, 0, 2, 0, 1, 2, 1, 2, 1, 0, 0, 0, 1, 0, 0, 2, 2,
           0, 2, 1, 2, 0, 0, 0, 0, 0, 0, 0, 2, 2, 2, 0, 0, 0, 2, 0, 0, 2, 0,
           1, 0, 0, 1, 2, 2, 2, 2, 2, 2, 0, 0, 1, 2, 0, 2, 1, 1, 2, 2, 0, 0,
           2, 0, 1, 0, 0, 2, 2, 0, 1, 1])
```

We now test the accuracy of our decision tree using the test set.

```
[39]: from sklearn.metrics import accuracy_score
```

```
[40]: accuracy_score(y_test, y_pred)
```

```
[40]: 0.9795918367346939
```

Visualise the decision tree

```
[41]: from sklearn.tree import plot_tree
```


```
[42]: plt.figure(figsize=(12, 12)) # set plot size (denoted in inches)
      plot_tree(clf,
                feature_names=list(X.columns),
                class_names=['Adelie', 'Chinstrap', 'Gentoo'])
```

```
[42]: [Text(0.6527777777777778, 0.9166666666666666, 'flipper_length_mm <= 206.5\ngini
= 0.639\nsamples = 226\nvalue = [101, 48, 77]\nclass = Adelie'),
      Text(0.4722222222222222, 0.75, 'bill_length_mm <= 44.65\ngini = 0.43\nsamples =
144\nvalue = [99, 45, 0]\nclass = Adelie'),
      Text(0.2777777777777778, 0.5833333333333334, 'bill_length_mm <= 42.35\ngini =
0.077\nsamples = 100\nvalue = [96, 4, 0]\nclass = Adelie'),
      Text(0.16666666666666666, 0.4166666666666667, 'bill_depth_mm <= 16.7\ngini =
0.022\nsamples = 92\nvalue = [91, 1, 0]\nclass = Adelie'),
      Text(0.11111111111111111, 0.25, 'bill_length_mm <= 39.3\ngini = 0.219\nsamples =
8\nvalue = [7, 1, 0]\nclass = Adelie'),
      Text(0.05555555555555555, 0.08333333333333333, 'gini = 0.0\nsamples = 7\nvalue
= [7, 0, 0]\nclass = Adelie'),
      Text(0.16666666666666666, 0.08333333333333333, 'gini = 0.0\nsamples = 1\nvalue
= [0, 1, 0]\nclass = Chinstrap'),
      Text(0.2222222222222222, 0.25, 'gini = 0.0\nsamples = 84\nvalue = [84, 0,
0]\nclass = Adelie'),
      Text(0.3888888888888889, 0.4166666666666667, 'sex <= 0.5\ngini = 0.469\nsamples
= 8\nvalue = [5, 3, 0]\nclass = Adelie'),
```

```

Text(0.3333333333333333, 0.25, 'gini = 0.0\nsamples = 3\nvalue = [0, 3,
0]\nclasse = Chinstrap'),
Text(0.4444444444444444, 0.25, 'gini = 0.0\nsamples = 5\nvalue = [5, 0,
0]\nclasse = Adelie'),
Text(0.6666666666666666, 0.5833333333333334, 'island_Torgersen <= 0.5\ngini =
0.127\nsamples = 44\nvalue = [3, 41, 0]\nclasse = Chinstrap'),
Text(0.6111111111111112, 0.4166666666666667, 'bill_depth_mm <= 20.75\ngini =
0.046\nsamples = 42\nvalue = [1, 41, 0]\nclasse = Chinstrap'),
Text(0.5555555555555556, 0.25, 'gini = 0.0\nsamples = 41\nvalue = [0, 41,
0]\nclasse = Chinstrap'),
Text(0.6666666666666666, 0.25, 'gini = 0.0\nsamples = 1\nvalue = [1, 0,
0]\nclasse = Adelie'),
Text(0.7222222222222222, 0.4166666666666667, 'gini = 0.0\nsamples = 2\nvalue =
[2, 0, 0]\nclasse = Adelie'),
Text(0.8333333333333334, 0.75, 'bill_depth_mm <= 17.65\ngini = 0.116\nsamples =
82\nvalue = [2, 3, 77]\nclasse = Gentoo'),
Text(0.7777777777777778, 0.5833333333333334, 'gini = 0.0\nsamples = 77\nvalue =
[0, 0, 77]\nclasse = Gentoo'),
Text(0.8888888888888888, 0.5833333333333334, 'bill_length_mm <= 46.55\ngini =
0.48\nsamples = 5\nvalue = [2, 3, 0]\nclasse = Chinstrap'),
Text(0.8333333333333334, 0.4166666666666667, 'gini = 0.0\nsamples = 2\nvalue =
[2, 0, 0]\nclasse = Adelie'),
Text(0.9444444444444444, 0.4166666666666667, 'gini = 0.0\nsamples = 3\nvalue =
[0, 3, 0]\nclasse = Chinstrap')]

```



We dropped 11 rows with missing values in the sex column, but it turns out that this column wasn't used. Maybe we could just drop the whole column and keep the rows that had data in the other columns. This would give us a little more data and might help to improve accuracy further.

1.3.1 Gini impurity

```
[43]: from collections import Counter

def gini_impurity(values):
    """
    Calculate the Gini impurity for a list of values.

    Args:
```

values (list): A list of categorical values.

Returns:

float: The Gini impurity of the list.

"""

```
counts = Counter(values)
total = len(values)
impurity = 1.0

for label in counts:
    prob_of_label = counts[label] / total
    impurity -= prob_of_label ** 2
return impurity
```

```
[44]: labels = ['Gentoo', 'Gentoo', 'Adelie', 'Chinstrap']
      #labels = ['Gentoo', 'Gentoo', 'Adelie', 'Adelie']
      #labels = ['Gentoo', 'Gentoo', 'Gentoo', 'Gentoo']

      gini_impurity(labels)
```

[44]: 0.625